

人工智能复习笔记

制作人：匡乾 Sun Yat-sen University School of Data and Computer Science

Search

formalization (形式化)

Property of Search 搜索的属性

Uninformed Search 无信息搜索

Breadth first 宽度优先

Depth first 深度优先

Uniform cost 一致代价搜索

Depth-limited search 深度受限搜索

Iterative deepening search 迭代加深搜索

Bidirectional search 双向搜索

无信息搜索总结

path checking / cycle checking 路径检测/环检测

路径检测

环检测

总结

Heuristic search 启发式搜索

Greedy best-first search (Greedy BFS) 最佳优先搜索

A* Search A*搜索

Admissible 可接纳性

Consistency (Monotonicity) 一致性、单调性

IDA* 迭代加深A*算法

Game tree search 博弈树搜索

basic definition

MiniMax Strategy

Alpha-beta pruning

Alpha cut

Beta cut

总结

CSP (Constraint satisfaction problem)约束满足问题

Formalization 形式化

backtracking 回溯算法

Forward checking 向前检测

MRV (Minimum Remaining Values Heuristics) 最小剩余启发式

GAC (Generalized Arc Consistency) 整体边一致

KRR(Knowledge representation and reasoning) 知识表示与推理

FOL(First-order logic) 一阶逻辑

Clausal form

Refutation

Converting first-order formulas into clausal form

Unification

Resolution

Answer extraction

Reasoning under Uncertainty 不确定推理

Probability in General

Bayesian Networks

graph + tables

Construct a Bayes Net

Inference

Variable Elimination

VE Algorithm:

D-Separation

Machine Learning 机器学习

Decision tree 决策树

Entropy 熵

Information gain 信息增益

Overfit 过度拟合

Bayes Learning 贝叶斯学习

Maximum a posteriori (极大后验MAP)

Maximum Likelihood (极大似然ML)

Artificial Neural Networks 神经网络

Linear regression 线性回归

Logistic regression 逻辑回归

Forward and backward phases

Search

- Problem solving by search: formalization
- Uninformed search: Breadth-First, Uniform-Cost, Depth-First, Depth-Limited, and Iterative- Deepening
- Heuristic search: Greedy best-first, A*
- Properties of search: completeness, optimality, time and space complexity
- Path/cycle checking
- Game tree search: MiniMax, alpha-beta pruning
- CSP: Formalization, backtracking, forward checking, and GAC algorithms

formalization (形式化)

1. Formulate a **state space** (形式化状态空间)
抽象真实问题
2. Formulate **actions** (形式化动作)

allow one to move between different states

3. Identify the **initial state** (确定初始状态)
4. Identify the **goal** or **desired condition** (确定目标)
5. Formulate heuristic (形式化启发式)

Example:

- States: the various cities you could be located in.
- Actions: drive between neighboring cities.
- Initial state: in Arad
- Goal: in Bucharest
- Solution: the route, the sequence of cities to travel through to get to Bucharest.

Property of Search 搜索的属性

- **Completeness 完备性**: will the search always find a solution if a solution exists?
- **Optimality 最优性**: will the search always find the least cost solution? (when actions have costs)
- **Time complexity 时间复杂度**: what is the maximum number of nodes that can be expanded or generated?
- **Space complexity 空间复杂度**: what is the maximum number of nodes that have to be stored in memory?

Uninformed Search 无信息搜索

Breadth first 宽度优先

将继承者放置到边界末端

example:

$\{0<>\}$

$\{1,2\}$

$\{2,2,3\}$

$\{2,3,3,4\}$

$\{3,3,4,3,4\}$

$\{3,4,3,4,4,5\}$

完备性、最优性: Yes

从小到大寻求方案，直到找到答案为止

最大继承数: b

最小解决方案步数: d

时间复杂度: $1 + b + b^2 + \dots + b^d + b(b^d - 1) = O(b^{d+1})$

空间复杂度: $b(b^d - 1) = O(b^{d+1})$

Depth first 深度优先

将继承者放置到边界前端

example:

$\{0\}$

$\{1,2\}$

$\{2,3,2\}$

$\{3,4,3,2\}$

$\{4,5,4,3,2\}$

$\{5,6,4,5,4,3,2\}$

完备性:

- Infinite state space: No
- Finite state space with infinite paths: No
- Finite state space and prune paths with duplicate states ? Yes

最优性: No

最大继承数: b

最小解决方案步数: d

时间复杂度: $O(b^m)$ m 是状态空间中最长的路径; 若 m 远远大于 d 则非常糟糕, 但若有多解往往会比较快

空间复杂度: $O(bm)$ 线性, 每次仅探索一条路径

Uniform cost 一致代价搜索

边界顺序由代价(cost)决定, 永远扩展代价最小的路径

完备性、最优性: Yes

C^* : 最优结果的代价 ϵ : 每一步的代价

时间、空间复杂度: $O(b^{C^*/\epsilon+1})$

Depth-limited search 深度受限搜索

设置的深度: L

- Completeness: No
- Optimality: No
- Time complexity: $O(b^L)$
- Space complexity: $O(bL)$

Iterative deepening search 迭代加深搜索

初始令 $L=0$, 并逐渐增大 L

- Completeness: Yes
- Optimality: Yes if costs are uniform

时间复杂度: $O(b^d)$

空间复杂度: $O(bd)$

Bidirectional search 双向搜索

Completeness: Yes

Optimality: if edges have uniform costs

Time and space complexity: $O(b^{d/2})$

无信息搜索总结

| Criterion | Breadth-First | Uniform-Cost | Depth-First | Depth-Limited | Iterative Deepening | Bidirectional (if applicable) |
|-----------|------------------|---------------------------------------|-------------|---------------|---------------------|-------------------------------|
| Complete? | Yes ^a | Yes ^{a,b} | No | No | Yes ^a | Yes ^{a,d} |
| Time | $O(b^d)$ | $O(b^{1+\lceil C^*/\epsilon \rceil})$ | $O(b^m)$ | $O(b^\ell)$ | $O(b^d)$ | $O(b^{d/2})$ |
| Space | $O(b^d)$ | $O(b^{1+\lceil C^*/\epsilon \rceil})$ | $O(bm)$ | $O(b\ell)$ | $O(bd)$ | $O(b^{d/2})$ |
| Optimal? | Yes ^c | Yes | No | No | Yes ^c | Yes ^{c,d} |

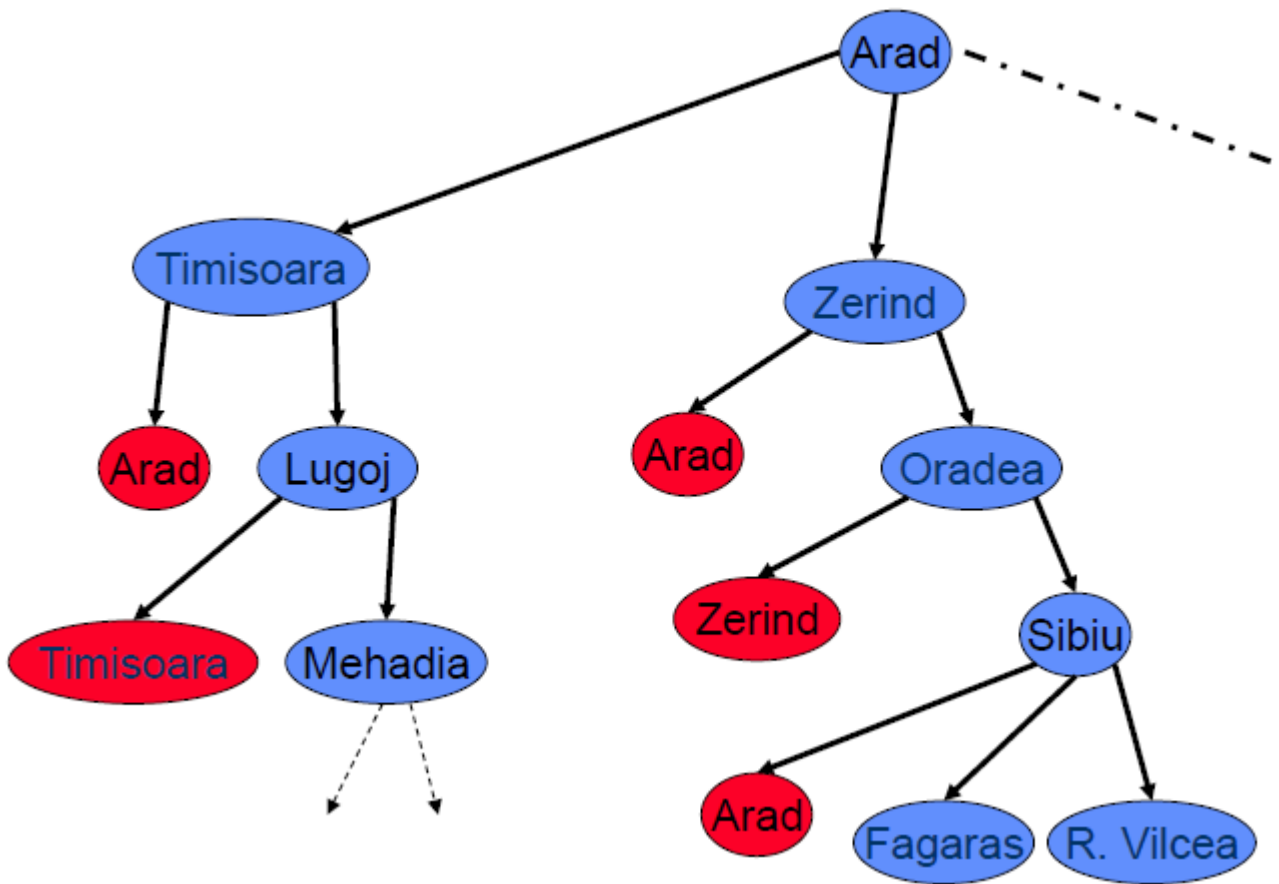
(BFS中的空间和时间改为 $O(b^{d+1})$)

path checking / cycle checking 路径检测/环检测

路径检测


通向c的路径: $\langle n_1, \dots, n_k, c \rangle$

则c不能与 n_i 相同



环检测

在整个探索过程中记录结点，确保扩展的结点c不与之前任何状态中的结点相同



总结

- Path checking: when we expand n to obtain child c , ensures that the state c is not equal to the state reached by any ancestor of c along this path
- Cycle checking: keep track of all states previously expanded during the search; when we expand n to obtain child c , ensure that c is not equal to any previously expanded state
- For uniform-cost search, cycle checking preserves optimality


Heuristic search 启发式搜索

idea: 得到启发式函数 $h(n)$, 预测从当前节点 n 到目标节点的花费

Greedy best-first search (Greedy BFS) 最佳优先搜索

用 $h(n)$ 对边界中的结点进行排序, 优先获取 low cost 的解

该方法忽略了到达 n 的 cost



Thus Greedy BFS is incomplete, not optimal

A* Search A*搜索

evaluation function 评估函数: $f(n) = g(n) + h(n)$

$g(n)$ 是到达结点 n 的路径花费

$h(n)$ 是启发式估计从结点 n 到达终点的花费

$f(n)$ 是对经过结点 n 到达终点的估计

Admissible 可接纳性

$h^*(n)$ 是从 n 到达终点的最佳路径的花费

$h(n)$ 是可容许的如果对于所有节点 n 都有 $h(n) \leq h^*(n)$

Admissible 可容纳的启发式低估真正的花费

$h(g) = 0$, 如果 n 不能到达终点则 $h(n) = \infty$

可接纳性 \rightarrow 最佳性 Admissibility implies optimality

Consistency (Monotonicity) 一致性、单调性

$h(n)$ 一致的/单调的, 如果对于任意结点 n_1, n_2 都有 $h(n_1) \leq c(n_1 \rightarrow n_2) + h(n_2)$

一致性 \rightarrow 可接纳性 Consistency implies admissibility

Note that consistency implies admissibility (proof)

- Case 1: no path from n to the goal
- Case 2: Let $n = n_1 \rightarrow n_2 \rightarrow \dots \rightarrow n_k$ be an optimal path from n to a goal node. We prove by induction on i that for all i , $h(n_i) \leq h^*(n_i)$.

单调性保证能在第一次到达某结点就是最佳路径

若没有单调性, 则需要记住之前路径的花费

性质:

1. **Proposition 1.** The f -values of nodes along a path must be non-decreasing

$f(n)$ 单调递增

2. **Proposition 2.** If n_2 is expanded after n_1 , then $f(n_1) \leq f(n_2)$

若 n_2 在 n_1 后出现, 则 $f(n_1) \leq f(n_2)$

3. **Proposition 3.** When n is expanded every path with lower f -value has already been expanded.

任何 f 花费小于 $f(n)$ 的结点必然已经扩展过

4. **Proposition 4.** The first time A^* expands a state, it has found the minimum cost path to that state.

第一次扩展到的结点就是最短路径

5. 在单调性的前提下, 换检测保证了最佳性

IDA* 迭代加深A*算法

迭代cutoff value为f-value, 而不是原来的L (深度)

边界中以f(n)的大小来排序

Theorem. The optimal cost to nodes in the relaxed problem is an admissible heuristic for the original problem!

放松问题中的最优花费是对于原问题可接受的启发式


Game tree search 博弈树搜索

basic definition

- Player: A(Max), B(Min)
- State: S
- Initial state: I
- Terminal state: T
- Successors
- Utility(效益), Payoff function: V

MiniMax Strategy

- $U(n) = \min \{U(c) : c \text{ is a child of } n\}$ if n is a Min node
- $U(n) = \max \{U(c) : c \text{ is a child of } n\}$ if n is a Max node



```

DFMiniMax(n, Player) //return Utility of state n given that
                        //Player is MIN or MAX

If n is TERMINAL
Return V(n) //Return terminal states utility
            //(V is specified as part of game)

//Apply Player's moves to get successor states.
ChildList = n.Successors(Player)
If Player == MIN
    return minimum of DFMiniMax(c, MAX) over c ∈ ChildList
Else //Player is MAX
    return maximum of DFMiniMax(c, MIN) over c ∈ ChildList

```

Alpha-beta pruning


Two types of pruning:

- pruning of max nodes (α -cuts)
- pruning of min nodes (β -cuts)

Alpha cut

- At a Max node n :
 - Let β be the lowest value of n 's siblings examined so far (siblings to the left of n that have already been searched)
 - Let α be the highest value of n 's children examined so far (changes as children examined)

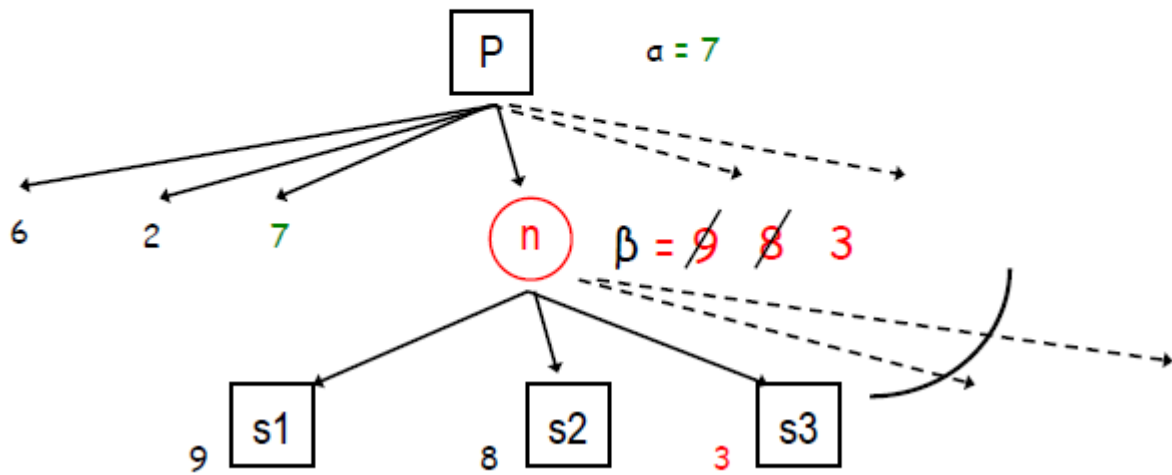
- While at a Max node n , if α becomes $\geq \beta$ we can stop expanding the children of n
- Min will never choose to move from n 's parent to n since it would choose one of n 's lower valued siblings first.



Beta cut

- At a Min node n :
 - Let α be the highest value of n 's sibling's examined so far (fixed when evaluating n)
 - Let β be the lowest value of n 's children examined so far (changes as children examined)

- If β becomes $\leq \alpha$ we can stop expanding the children of n .
- Max will never choose to move from n 's parent to n since it would choose one of n 's higher value siblings first.



总结

当 $\beta \leq \alpha$ 时, 进行剪枝

Minimax 需要探索 $O(b^D)$ 个结点, 而alpha-beta剪枝需要探索 $O(b^{D/2})$ 个结点

CSP (Constraint satisfaction problem)约束满足问题

Formalization 形式化

A CSP consists of:

- A set of variables: V_1, \dots, V_n
- Each variable has a domain: $\text{Dom}[V_i]$ ($V_i = d \iff d \in \text{Dom}[V_i]$)
- A set of constraints: C_1, \dots, C_m e.g. $C(V_1, V_2, V_4)$

goal: 寻找满足条件的解, 使得各个变量都有取值

backtracking 回溯算法

- We pick a variable*,
- pick a value for it*,
- test the constraints that we can,
- if a constraint is unsatisfied we backtrack,
- otherwise we set another variable.
- When all the variables are set, we're done.

启发式应用于挑选变量和挑选值：

the order in which variables are assigned:


PickUnassignedVariable()

the order of values tried for each variable.

Forward checking 向前检测

检查那些只含有一个未实例化变量的约束，去除那个变量所有违反约束取值

同时要记住，每一个值是在哪一步被去除的



MRV (Minimum Remaining Values Heuristics) 最小剩余启发式

先执行值域较小的变量，当一个变量只有一个取值时，立即执行

- What variables would you try first?

| | | | | | | | |
|---|---|---|---|---|---|---|---|
| 8 | 1 | 5 | 6 | | | | 4 |
| 6 | | | | 7 | 5 | | 8 |
| | | | | 9 | | | |
| 9 | | | | 4 | 1 | 7 | |
| | 4 | | | | | | 2 |
| | | 6 | 2 | 3 | | | 8 |
| | | | | 5 | | | |
| | 5 | | 9 | 1 | | | 6 |
| 1 | | | | | 7 | 8 | 9 |
| | | | | | | | 5 |

Domain of each variable:
 $\{1, \dots, 9\}$

(1, 5) impossible values:

Row: $\{1, 4, 5, 6, 8\}$

Column: $\{1, 3, 4, 5, 7, 9\}$

Subsquare: $\{5, 6, 7, 9\}$

→ Domain = $\{2\}$

(9, 5) impossible values:

Row: $\{1, 5, 7, 8, 9\}$

Column: $\{1, 3, 4, 5, 7, 9\}$

Subsquare: $\{1, 5, 7, 9\}$

→ Domain = $\{2, 6\}$

After assigning value 2 to
 cell (1,5): Domain = $\{6\}$

Most restricted variables! = MRV

GAC (Generalized Arc Consistency) 整体边一致

Some definition:

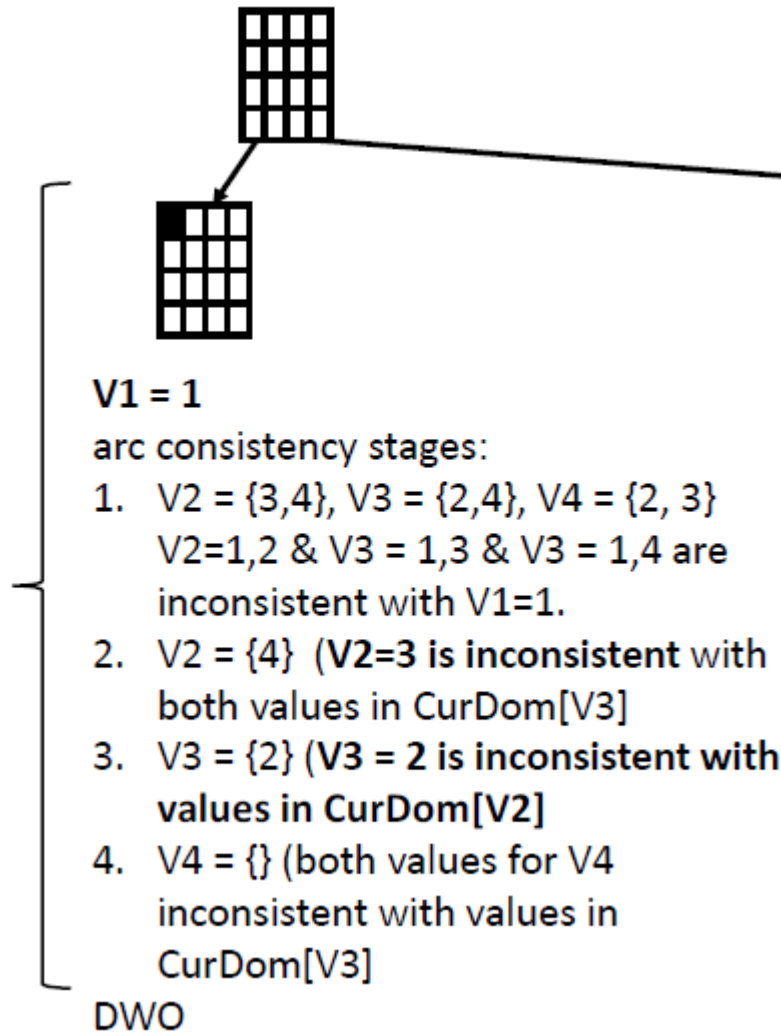
- $C(X,Y)$ is consistent $\iff \forall x, \exists y$ 满足 C
- $C(V_1, V_2, \dots, V_n)$ 关于 V_i is GAC $\iff \forall V_i, \exists V_1, \dots, V_{i-1}, V_{i+1}, \dots, V_n$ 满足 C
- A constraint(C) is GAC \iff 关于它的任何变量都是GAC的
- A CSP is GAC \iff 所有限制(C)都是GAC的

如果对于变量 V , 取值 d 不能得到一个解, 这说 d 是arc inconsistent (边不一致的)

$C(X,Y): X > Y, \text{Dom}(X)=\{1,5,11\}, \text{Dom}(Y)=\{3,8,15\}$

- For $X=1$ there is no value of Y s.t. $1 > Y$, so remove 1 from domain X
- For $Y=15$ there is no value of X s.t. $X > 15$, so remove 15 from domain Y
- We obtain $\text{Dom}(X)=\{5,11\}$ and $\text{Dom}(Y)=\{3,8\}$.

GAC检查的过程需要不断的循环, 因为一个定义域改变可能引起其它定义域变化



GAC必须在每个节点都检查所有限制(C)

Example: http://www.cs.toronto.edu/~fbacchus/csc384/Lectures/Tutorial3_CSP.pdf

KRR(Knowledge representation and reasoning) 知识表示与推理

- First-order logic: syntax and semantics
- Soundness and completeness of proof procedures
- Converting first-order formulas into clausal form
- Unification and MGU
- Resolution proof: forward chaining and refutation
- Answer extraction

知识表示假设：所有AI system都是基于知识的(knowledge-based)

FOL(First-order logic) 一阶逻辑

$$a \rightarrow b \iff \neg a \vee b$$

$$a \leftrightarrow b \iff (a \rightarrow b) \wedge (b \rightarrow a)$$

Clausal form

e.g., $p \vee \neg r \vee s$, written $(p, \neg r, s)$

Proposition. $\{p\} \cup c_1, \{\neg p\} \cup c_2 \models c_1 \cup c_2$

Refutation

$KB \models \alpha$ iff $KB \wedge \neg \alpha$ is unsatisfiable

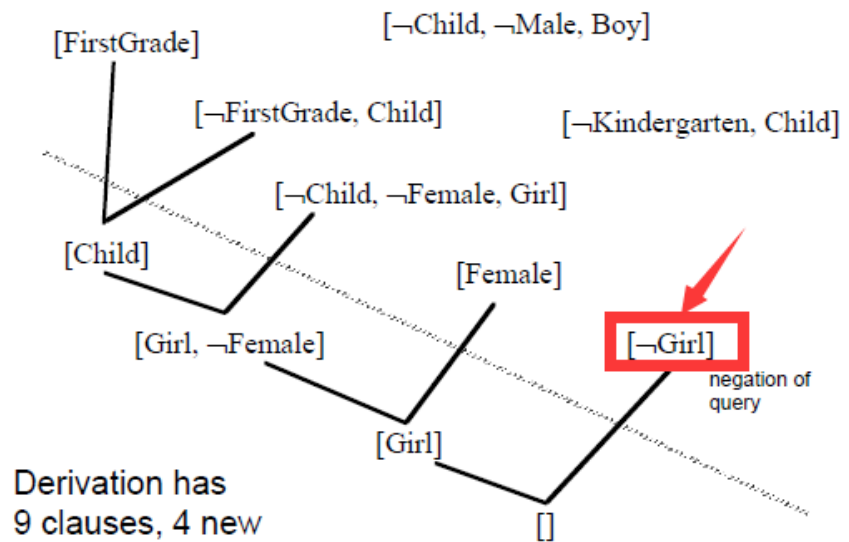
Thus to check if $KB \models \alpha$,

- put KB and $\neg \alpha$ into clausal form to get S,
- check if $S \vdash ()$

KB

FirstGrade
 FirstGrade \supset Child
 Child \wedge Male \supset Boy
 Kindergarten \supset Child
 Child \wedge Female \supset Girl
 Female

Show that **KB \models Girl**



Converting first-order formulas into clausal form

Step:

1. Eliminate Implications (消去蕴含)

$$A \rightarrow B \Leftrightarrow \neg A \vee B$$

2. Move negations inwards using (将括号外，量词外的非挪到里面)

$$\bullet \neg(A \vee B) \Leftrightarrow \neg A \wedge \neg B, \neg(A \wedge B) \Leftrightarrow \neg A \vee \neg B$$

$$\bullet \neg \exists x.A \Leftrightarrow \forall x.\neg A, \neg \forall x.A \Leftrightarrow \exists x.\neg A, \neg \neg A \Leftrightarrow A$$

3. Standardize Variables (规范变量名称，使每个量化变量都unique)

$$\forall x\{\neg P(x) \vee [\forall y(\neg P(y) \vee P(f(x, y))) \wedge \exists y(Q(x, y) \vee \neg P(y))]\}$$

3. Standardize Variables (Rename variables so that each quantified variable is unique)

$$\forall x\{\neg P(x) \vee [\forall y(\neg P(y) \vee P(f(x, y))) \wedge \exists z(Q(x, z) \vee \neg P(z))]\}$$

4. Skolemize (将所有带有存在量词的变量，转换为关于全称量词变量的函数)

$$\forall x \{ \neg P(x) \vee [\forall y (\neg P(y) \vee P(f(x, y))) \wedge \exists z (Q(x, z) \vee \neg P(z))] \}$$

4. Skolemize (Remove existential quantifiers by introducing new function symbols)

$$\forall x \{ \neg P(x) \vee [\forall y (\neg P(y) \vee P(f(x, y))) \wedge (Q(x, g(x)) \vee \neg P(g(x)))] \}$$

5. Convert to prenex form (转换为前束范式, 即将所有量词提到最前面)

6. Disjunctions over conjunctions (把交提出来)

$$A \vee (B \wedge C) \iff (A \vee B) \wedge (A \vee C)$$

7. Flatten nested conjunctions and disjunctions (不知道干嘛的)

8. Convert to Clauses (去除量词, 把交分开)

$$\forall x \forall y \{ (\neg P(x) \vee \neg P(y) \vee P(f(x, y))) \wedge (\neg P(x) \vee Q(x, g(x)) \vee \neg P(g(x))) \}$$

8. Convert to Clauses (remove quantifiers and break apart conjunctions).

$$\text{a) } \neg P(x) \vee \neg P(y) \vee P(f(x, y))$$

$$\text{b) } \neg P(x) \vee Q(x, g(x)) \vee \neg P(g(x))$$

Unification

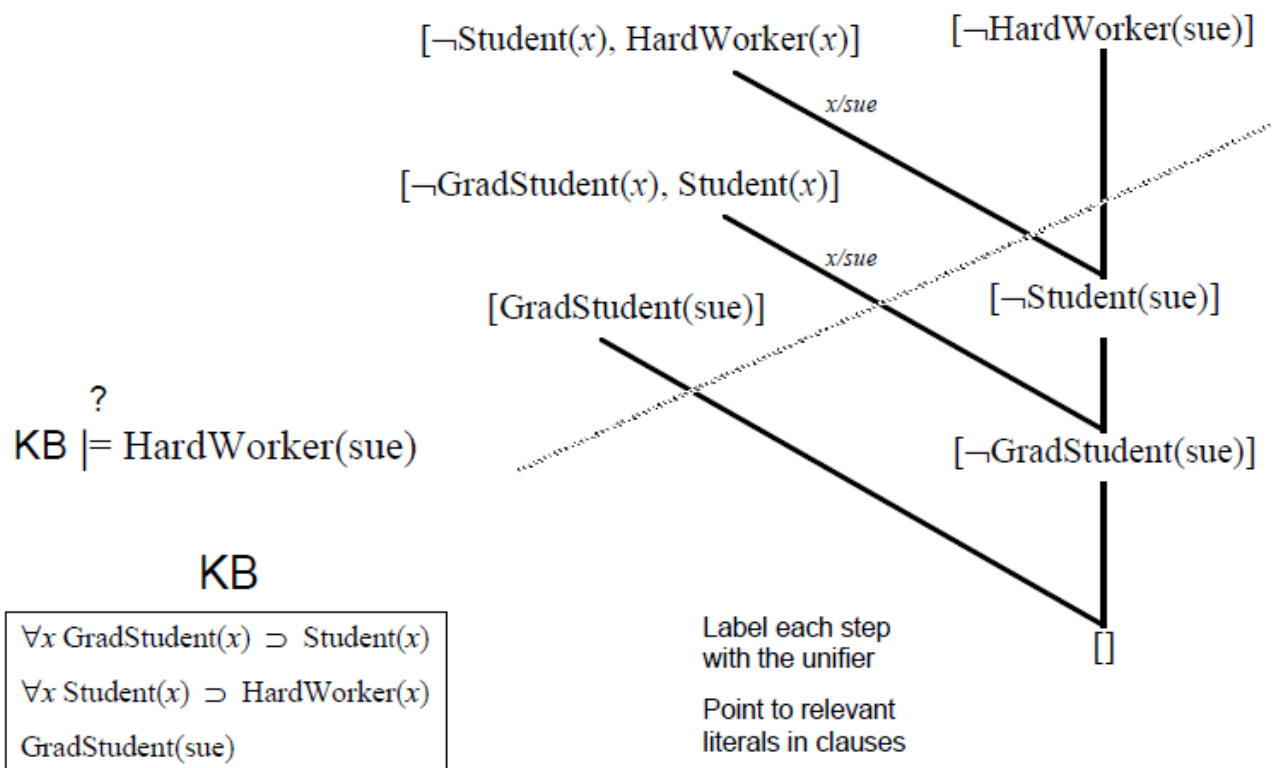
- Let $\theta = \{x = f(y), y = z\}$, $\sigma = \{x = a, y = b, z = y\}$
- Step 1. Get $S = \{x = f(b), y = y, x = a, y = b, z = y\}$
- Step 2. Delete $y = y$.
- Step 3. Delete $x = a$.
- The result is $S = \{x = f(b), y = b, z = y\}$

Resolution

example:

1. $(P(x), Q(g(x)))$
2. $(R(a), Q(z), \neg P(a))$
3. $R[1a, 2c]\{X=a\} (Q(g(a)), R(a), Q(z))$

- “R” means resolution step.



Prove that $\exists y \forall x P(x, y) \models \forall x \exists y P(x, y)$


- $\exists y \forall x P(x, y) \Rightarrow 1. P(x, a)$
- $\neg \forall x \exists y P(x, y) \Leftrightarrow \exists x \forall y \neg P(x, y) \Rightarrow 2. \neg P(b, y)$
- $R[1, 2]\{x = b, y = a\}()$

Answer extraction

- We can also answer wh- questions
- Replace query $\exists x P(x)$ by $\exists x [P(x) \wedge \neg \text{answer}(x)]$
- Instead of deriving (), derive any clause containing just the answer predicate

KB: Student(john)
 Student(jane)
 Happy(john)

Q: $\exists x [\text{Student}(x) \wedge \text{Happy}(x)]$



直接在Clausal form下的query插入answer(x)即可

Reasoning under Uncertainty 不确定推理

- Bayesian networks: graphs + tables, inference
- Variable elimination algorithm
- Use D-separation to determine independence

Probability in General

- $Pr(U) = 1$
- $Pr(A) \in [0, 1]$
- $Pr(A \cup B) = Pr(A) + Pr(B) - Pr(A \cap B)$

$$Pr(\{V_1 = a\}) = \sum_{x_2 \in D[V_2]} \cdots \sum_{x_n \in D[V_n]} Pr(V_1 = a, V_2 = x_2, \dots, V_n = x_n)$$

Conditional probabilities (条件概率) :

$$Pr(B|A) = Pr(B \cap A)/Pr(A)$$

全集分割:

$$B_1, B_2, \dots, B_k$$

(不交, 不漏)

- $B_i \cap B_j = \emptyset, i \neq j$ (mutually exclusive)
- $B_1 \cup B_2 \cup \dots \cup B_k = U$ (exhaustive)

In probabilities:

- $Pr(B_i \cap B_j) = 0, i \neq j$
- $Pr(B_1 \cup B_2 \cup \dots \cup B_k) = 1$

Sumout rule:

$$Pr(A) = Pr(A \cap B_1) + \dots + Pr(A \cap B_k)$$

In conditional probabilities:

$$Pr(A) = Pr(A|B_1)Pr(B_1) + \dots + Pr(A|B_k)Pr(B_k)$$

Independent:

$Pr(B|A) = Pr(B)$ (B is independent of A)

- If A and B are independent, then

$$Pr(A \cap B) = Pr(A) \cdot Pr(B)$$
- If given A , B and C are conditionally independent, then

$$Pr(B \cap C|A) = Pr(B|A) \cdot Pr(C|A)$$

Bayes rule:

$$Pr(Y|X) = Pr(X|Y)Pr(Y)/Pr(X)$$

Chain rule:

$$Pr(A_1 \cap A_2 \cap \dots \cap A_n) = Pr(A_1|A_2 \cap \dots \cap A_n) \cdot Pr(A_2|A_3 \cap \dots \cap A_n) \cdot \dots \cdot Pr(A_{n-1}|A_n) \cdot Pr(A_n)$$

Notation / Terminology:

$\Pr(X) == \Pr(X=d)$ for all d in $\text{Dom}[X]$

$$\sum_{d \in \text{Dom}[X]} \Pr(X = d) = 1$$

Inference:



- Computing $\Pr(a)$ in more concrete terms:

- $\Pr(c) = \Pr(c|e)\Pr(e) + \Pr(c|\sim e)\Pr(\sim e)$
 $= 0.9 * 0.7 + 0.5 * 0.3 = 0.78$
- $\Pr(\sim c) = \Pr(\sim c|e)\Pr(e) + \Pr(\sim c|\sim e)\Pr(\sim e) = 0.22$
 - $\Pr(\sim c) = 1 - \Pr(c)$, as well
- $\Pr(a) = \Pr(a|c)\Pr(c) + \Pr(a|\sim c)\Pr(\sim c)$
 $= 0.3 * 0.78 + 1.0 * 0.22 = 0.454$
- $\Pr(\sim a) = 1 - \Pr(a) = 0.546$

Bayesian Networks

graph + tables


A BN over variables $\{X_1, X_2, \dots, X_n\}$ consists of:

- a DAG (directed acyclic graph) whose nodes are the variables
- a set of CPTs (conditional probability tables)
 $\Pr(X_i | \text{Par}(X_i))$ for each X_i

example:

$\Pr(A,B,C,D,E,F,G,H,I,J,K) =$

$\Pr(A)$
 $\times \Pr(B)$
 $\times \Pr(C|A)$
 $\times \Pr(D|A,B)$
 $\times \Pr(E|C)$
 $\times \Pr(F|D)$
 $\times \Pr(G)$
 $\times \Pr(H|E,F)$
 $\times \Pr(I|F,G)$
 $\times \Pr(J|H,I)$
 $\times \Pr(K|I)$



Construct a Bayes Net

- Step 1 Apply the Chain Rule

$$\Pr(X_1, \dots, X_n) = \Pr(X_n | X_1, \dots, X_{n-1}) \Pr(X_{n-1} | X_1, \dots, X_{n-2}) \dots \Pr(X_1)$$

- Step 2 移除所有无关变量

$$\Pr(X_n | \text{Par}(X_n)) \Pr(X_{n-1} | \text{Par}(X_{n-1})) \dots \Pr(X_1)$$

- Step 3 建立一个图(DAG)
- Step 4 确定CPT(conditional probability table)条件概率表格

Inference

Given

1) a Bayes net

$$\begin{aligned} \Pr(X_1, X_2, \dots, X_n) \\ = \Pr(X_n \mid \text{Par}(X_n)) * \Pr(X_{n-1} \mid \text{Par}(X_{n-1})) * \dots * \Pr(X_1 \mid \text{Par}(X_1)) \end{aligned}$$

2) some Evidence, E

$E = \{\text{a set of values for some of the variables}\}$

We want to

- compute the new probability distribution

$$\Pr(X_k \mid E)$$

That is, we want to figure out

$$\Pr(X_k = d \mid E) \text{ for all } d \in \text{Dom}[X_k]$$

Variable Elimination

Variable elimination uses

- the product decomposition, and
- the summing out rule
- In general, at each stage VE will compute a table of numbers: one for each different instantiation of the variables in the sum.

- Let $f(\underline{X}, \underline{Y})$ & $g(\underline{Y}, \underline{Z})$ be two factors with variables \underline{Y} in common
- The **product** of f and g , denoted $h = f \times g$ (or sometimes just $h = fg$), is defined:

$$h(\underline{X}, \underline{Y}, \underline{Z}) = f(\underline{X}, \underline{Y}) \times g(\underline{Y}, \underline{Z})$$

| f(A,B) | | g(B,C) | | h(A,B,C) | | | |
|--------|-----|--------|-----|----------|------|--------|------|
| ab | 0.9 | bc | 0.7 | abc | 0.63 | ab~c | 0.27 |
| a~b | 0.1 | b~c | 0.3 | a~bc | 0.08 | a~b~c | 0.02 |
| ~ab | 0.4 | ~bc | 0.8 | ~abc | 0.28 | ~ab~c | 0.12 |
| ~a~b | 0.6 | ~b~c | 0.2 | ~a~bc | 0.48 | ~a~b~c | 0.12 |

restrict a Factor:

- Let $f(X, \underline{Y})$ be a factor with variable X (\underline{Y} is a set)
- We **restrict** factor f to $X=a$ by setting X to the value a and “deleting” incompatible elements of f ’s domain .
Define $h = f_{X=a}$ as: $h(\underline{Y}) = f(a, \underline{Y})$

| f(A,B) | | h(B) = f _{A=a} | |
|--------|-----|-------------------------|-----|
| ab | 0.9 | b | 0.9 |
| a~b | 0.1 | ~b | 0.1 |
| ~ab | 0.4 | | |
| ~a~b | 0.6 | | |

VE Algorithm:

Given:


- Bayes Net with CPTs factors F ,
- query variable Q ,
- evidence variables E (observed to have values e),
- remaining variables Z .

Now Compute $\Pr(Q|E)$

- 1 Replace each factor $f \in F$ that mentions a variable(s) in E with its restriction $f_{E=e}$ (this might yield a “constant” factor)
- 2 For each Z_j — in the order given —eliminate $Z_j \in Z$ as follows:
 - 1 Let f_1, f_2, \dots, f_k be the factors in F that include Z_j
 - 2 Compute new factor $g_j = \sum_{Z_j} f_1 \times f_2 \times \dots \times f_k$
 - 3 Remove the factors f_i from F and add new factor g_j to F
- 3 The remaining factors refer only to the query variable Q . Take their product and normalize to produce $\Pr(Q|E)$.

1. 用已知事实替换变量
2. 将变量 Z_j 用关于其它变量的函数表示, 从而消去 Z_j
将包含 Z_j 的用 f_i 表示, 并将它们全部消去最后加入新产生的 g_i
3. 最后只剩下查询变量

Factors: $f_1(A)$ $f_2(B)$ $f_3(A,B,C)$
 $f_4(C,D)$
Query: $P(A)?$
Evidence: $D = d$
Elim. Order: C, B




Restriction: replace $f_4(C,D)$ with $f_5(C) = f_4(C,d)$

Step 1: **Eliminating C:** Compute & Add $f_6(A,B) = \sum_C f_5(C) f_3(A,B,C)$
 Remove: $f_3(A,B,C)$, $f_5(C)$

Step 2: **Eliminating B:** Compute & Add $f_7(A) = \sum_B f_6(A,B) f_2(B)$
 Remove: $f_6(A,B)$, $f_2(B)$

Last factors: $f_7(A)$, $f_1(A)$. The product $f_1(A) \times f_7(A)$ is (unnormalized) posterior. So... $P(A|d) = \alpha f_1(A) \times f_7(A)$
 where $\alpha = 1/\sum_A f_1(A)f_7(A)$ ← ****Note the Normalization Constant!****


D-Separation



If Z is in evidence, the path between X and Y is blocked

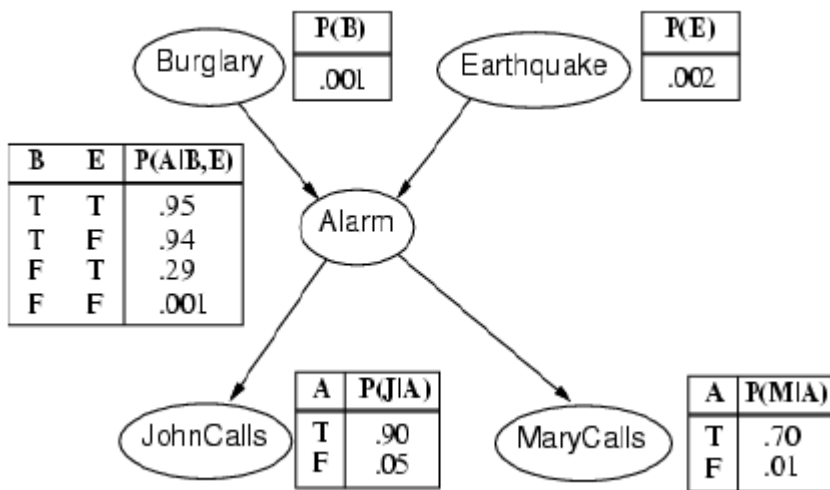


If Z is in evidence, the path between X and Y is blocked



If Z is **not** in evidence and **no** descendant of Z is in evidence, then the path between X and Y is blocked

example:



- A and M are dependent given J
- B and M are independent, given A
- J and M are dependent, but independent given A
- B and E are independent
- B and E are dependent, given A, J, or M

Machine Learning 机器学习

What is ML?

performance in T measured by P improves with E
(experience E, task T, performance measure P)

- Decision-tree learning
- Naive Bayes learning
- K-means and EM
- Chain rule for computing partial derivatives
- Linear and logistic regression
- Backpropagation
- Q-learning

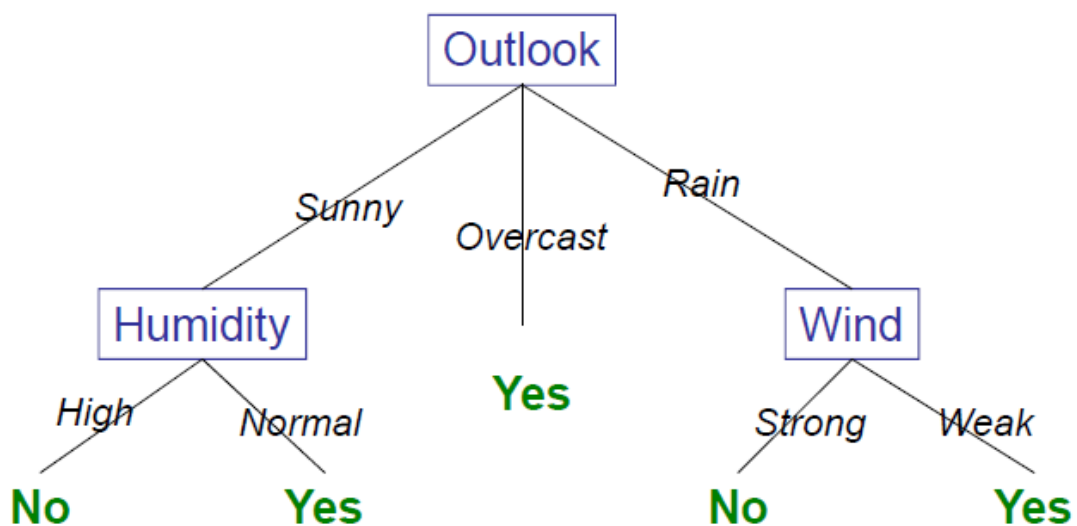
Decision tree 决策树

结点：标识属性

边：标识属性值

叶子：标识输出值

example:



An instance

<Outlook=Sunny, Temp=Hot, Humidity=High, Wind=Strong>

Classification: No

为了构造一个尽量小的树，我们应该优先选择具有代表性的属性

```
function DECISION-TREE-LEARNING(examples, attributes, parent_examples) returns  
a tree  
  
  if examples is empty then return PLURALITY-VALUE(parent_examples)  
  else if all examples have the same classification then return the classification  
  else if attributes is empty then return PLURALITY-VALUE(examples)  
  else  
     $A \leftarrow \operatorname{argmax}_{a \in \text{attributes}} \text{IMPORTANCE}(a, \text{examples})$   
    tree  $\leftarrow$  a new decision tree with root test A  
    for each value  $v_k$  of A do  
       $\text{exs} \leftarrow \{e : e \in \text{examples} \text{ and } e.A = v_k\}$   
      subtree  $\leftarrow$  DECISION-TREE-LEARNING(exs, attributes - A, examples)  
      add a branch to tree with label (A =  $v_k$ ) and subtree subtree  
  return tree
```

Entropy 熵

$$H(V) = - \sum_k P(v_k) \log_2 P(V_k)$$

The entropy of a Boolean random variable that is true with probability q :

$$B(q) = -(q \log_2 q + (1 - q) \log_2 (1 - q))$$

如果训练集有 p 个正确 n 个不正确的例子，则熵为：

$$H(\text{Goal}) = B\left(\frac{p}{p + n}\right)$$

Information gain 信息增益

So the expected entropy remaining after testing attribute A is

$$\text{Remainder}(A) = \sum_{k=1}^d \frac{p_k + n_k}{p + n} B\left(\frac{p_k}{p_k + n_k}\right).$$

p_k / n_k : positive/negative examples of the subset

The information gain (IG) from the attribute test on A is the expected reduction in entropy:

$$Gain(A) = B\left(\frac{p}{p+n}\right) - Remainder(A)$$

example:



- For the training set, $p = n = 6$, $B(6/12) = 1$
- $Gain(Pat) = 1 - \left[\frac{2}{12}B\left(\frac{0}{2}\right) + \frac{4}{12}B\left(\frac{4}{4}\right) + \frac{6}{12}B\left(\frac{2}{6}\right) \right] \approx 0.541$
- $Gain(Type) = 1 - \left[\frac{2}{12}B\left(\frac{1}{2}\right) + \frac{2}{12}B\left(\frac{1}{2}\right) + \frac{4}{12}B\left(\frac{2}{4}\right) + \frac{4}{12}B\left(\frac{2}{4}\right) \right] = 0$

Overfit 过度拟合

可以通过剪枝来去除关联度小的结点从而避免过度拟合

Bayes Learning 贝叶斯学习

- Prior: $Pr(H)$
- Likelihood: $Pr(d|H)$
- Evidence: $d = \langle d_1, d_2, \dots, d_n \rangle$
- Computing the posterior using Bayes' Theorem:

$$Pr(H|d) = \alpha Pr(d|H) Pr(H)$$

$$P(X|d) = \sum_i P(X|d, h_i) P(h_i|d) = \sum_i P(X|h_i) P(h_i|d)$$

Maximum a posteriori (极大后验MAP)

- Idea: make prediction based on most probable hypothesis
 - $h_{\text{MAP}} = \operatorname{argmax}_{h_i} P(h_i|d)$
 - $P(X|d) \approx P(X|h_{\text{MAP}})$

$$h_{\text{MAP}} = \operatorname{argmax}_h P(h) P(d|h)$$

需要考虑各个糖果方案出现的可能性即可得到 h_{MAP}

Maximum Likelihood (极大似然ML)

$$h_{\text{ML}} = \operatorname{argmax}_h P(d|h)$$

$$P(X|d) \approx P(X|h_{\text{ML}})$$

无需考虑各个糖果方案出现的可能性即可得到 h_{ML}

example:

- Hypothesis h_θ
 - $P(cherry) = \theta$ and $P(lime) = 1 - \theta$
- Data d :
 - c cherries and l limes

| $P(F=cherry)$ |
|---------------|
| θ |

Flavor

- $c/\theta - l/(1 - \theta) = 0 \Rightarrow \theta = c/(c + l)$

Artificial Neural Networks 神经网络

Loss function:

- Three commonly used loss functions:
 - Absolute value loss: $L_1(y, y') = |y - y'|$
 - Squared error loss: $L_2(y, y') = (y - y')^2$
 - 0/1 loss: $L_{0/1}(y, y') = 0$ if $y = y'$, else 1

Linear regression 线性回归

梯度下降:


$$w_i \leftarrow w_i - \alpha \partial \text{Loss}(w) / \partial w_i$$

α : 学习率

- $h_w(x) = w \cdot x = \sum_i w_i x_i$
- Squared error loss: $\text{Loss}(w) = (y - h_w(x))^2$
- Chain rule: $\partial g(f(x)) / \partial x = g'(f(x)) \partial f(x) / \partial x$
- $\partial \text{Loss}(w) / \partial w_i = -2(y - h_w(x))x_i$
- $w_i \leftarrow w_i + \alpha(y - h_w(x))x_i$

Logistic regression 逻辑回归

Sigmoid



$$g(x) = 1/(1+e^{-x})$$

logistic function:

$$g(x) = 1/(1 + e^{-x})$$

- $g(x) = 1/(1 + e^{-x})$
- $h_w(x) = g(w \cdot x)$
- $g' = g(1 - g)$
- $Loss(w) = (y - h_w(x))^2$
- $\begin{aligned} \partial Loss(w)/\partial w_i &= -2(y - h_w(x))g'(w \cdot x)x_i \\ &= -2(y - h_w(x))h_w(x)(1 - h_w(x))x_i \end{aligned}$
- $w_i \leftarrow w_i + \alpha(y - h_w(x))h_w(x)(1 - h_w(x))x_i$

```

initialize  $w$  arbitrarily
repeat
  for each  $e$  in examples do
     $p \leftarrow g(w \cdot x(e))$ 
     $\delta \leftarrow y(e) - p$ 
    for each  $i$  do
       $w_i \leftarrow w_i + \alpha \delta p(1 - p)x_i$ 
until some stopping criterion is satisfied
return  $w$ 

```

Forward and backward phases

Forward phase:

Output a_j at unit j : $a_j = g(in_j)$ where $in_j = \sum_i w_{ij}a_i$

$$in = \sum_i w_{ij}a_i \quad out = g(in) = 1/(1+e^{-in})$$

Backward phase:

- For an output unit j :

$$\Delta_j = g'(in_j)(y_j - a_j) = a_j(1 - a_j)(y_j - a_j)$$

- For an hidden unit i :

$$\Delta_i = g'(in_i) \sum_j w_{ij} \Delta_j = a_i(1 - a_i) \sum_j w_{ij} \Delta_j$$

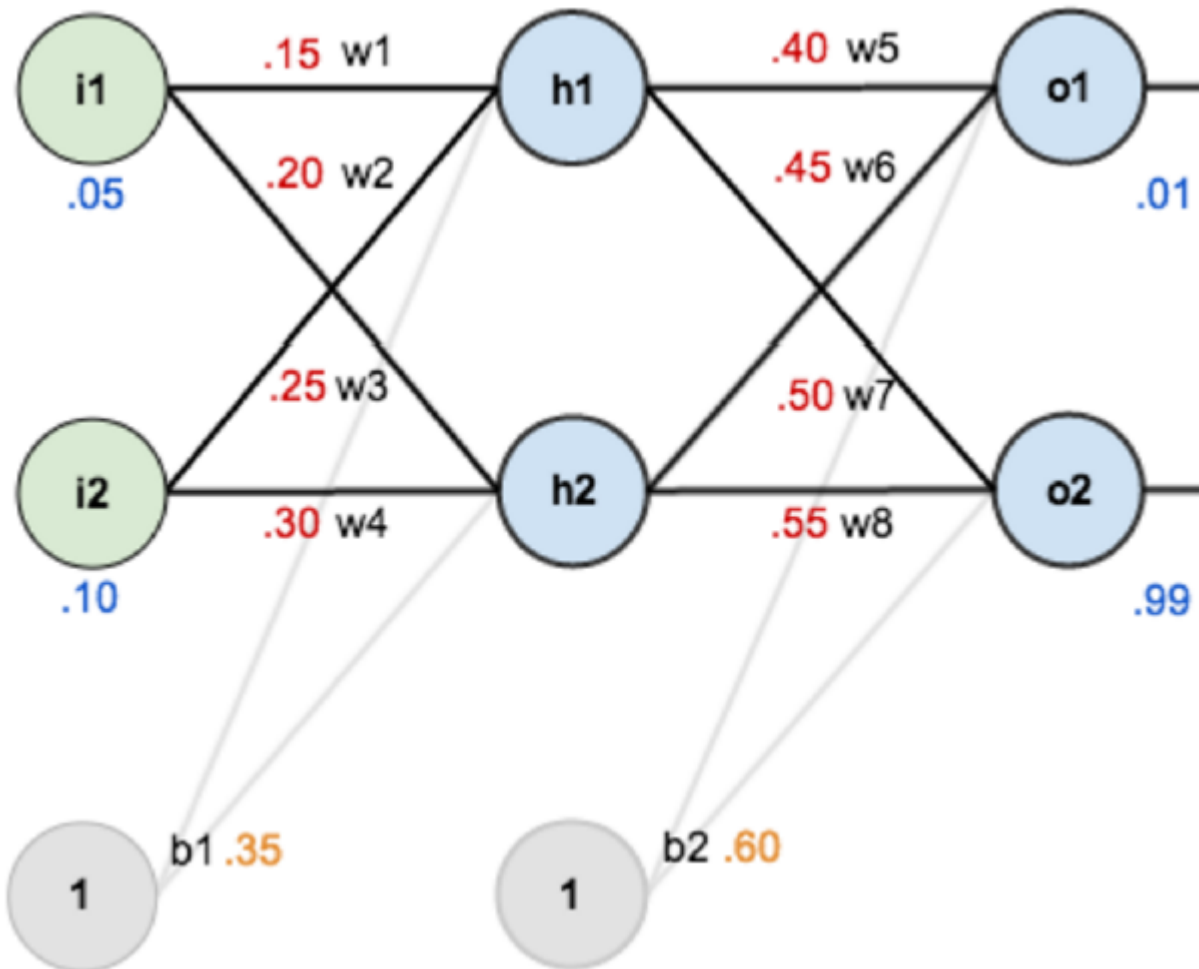
Weight updating: $w_{ij} \leftarrow w_{ij} + \alpha a_i \delta_j$

$$\Delta_o = o(1 - o)(y - o)$$

$$w^+ = w - \alpha * out * \Delta_o$$

$$\Delta h = out(1 - out) \sum_i w_i \Delta_o$$

example:



- $in_{h_1} = w_1 i_1 + w_2 i_2 + b_1 = 0.05 * 0.15 + 0.10 * 0.20 + 0.35 = 0.3775$
- $out_{h_1} = g(in_{h_1}) = \frac{1}{1+e^{-0.3775}} = 0.593269992$
- $out_{h_2} = 0.596884378$
- $in_{o_1} = w_5 out_{h_1} + w_6 out_{h_2} + b_2 = 0.40 * 0.593269992 + 0.45 * 0.596884378 + 0.60 = 1.105905967$
- $out_{o_1} = g(in_{o_1}) = \frac{1}{1+e^{-1.105905967}} = 0.75136507$
- $out_{o_2} = 0.772928465$

Let $\alpha = 0.5$

- $\Delta_{o_1} = 0.75136507(1 - 0.75136507)(0.01 - 0.75136507) = -0.138498562$

- $w_5^+ = w_5 + \alpha \cdot out_{h_1} \cdot \Delta_{o_1} =$
 $0.40 - 0.5 * 0.593269992 * 0.138498562 = 0.35891648$
- $w_6^+ = w_6 + \alpha \cdot out_{h_2} \cdot \Delta_{o_1} =$
 $0.45 - 0.5 * 0.596884378 * 0.138498562 = 0.408666186$
- $\Delta_{o_2} = 0.772928465(1 - 0.772928465)(0.99 - 0.772928465) =$
 0.0380982366
- $w_7^+ = w_7 + \alpha \cdot out_{h_1} \cdot \Delta_{o_2} =$
 $0.50 + 0.5 * 0.593269992 * 0.0380982366 = 0.511301270$
- $w_8^+ = w_8 + \alpha \cdot out_{h_2} \cdot \Delta_{o_2} =$
 $0.55 + 0.5 * 0.596884378 * 0.0380982366 = 0.561370121$
- $\Delta_{h_1} = g'(in_{h_1})(w_5\Delta_{o_1} + w_7\Delta_{o_2}) =$
 $0.593269992(1 - 0.593269992)(0.40 * (-0.138498562) +$
 $0.50 * 0.0380982366) = -0.241300709 * 0.036350306$
- $w_1^+ = w_1 + \alpha \cdot i_1 \cdot \Delta_{h_1} =$
 $0.15 - 0.5 * 0.05 * 0.241300709 * 0.036350306 = 0.149780716$
- $w_2^+ = 0.19956143$
- $w_3^+ = 0.24975114$
- $w_4^+ = 0.29950229$